


.
.

Artificial Intelligence (AI) for Engineering

COS40007

Dr. Afzal Azeem Chowdhary
Lecturer, SoCET, Swinburne University of Technology

Seminar 9: 29th April 2025



. .
. .

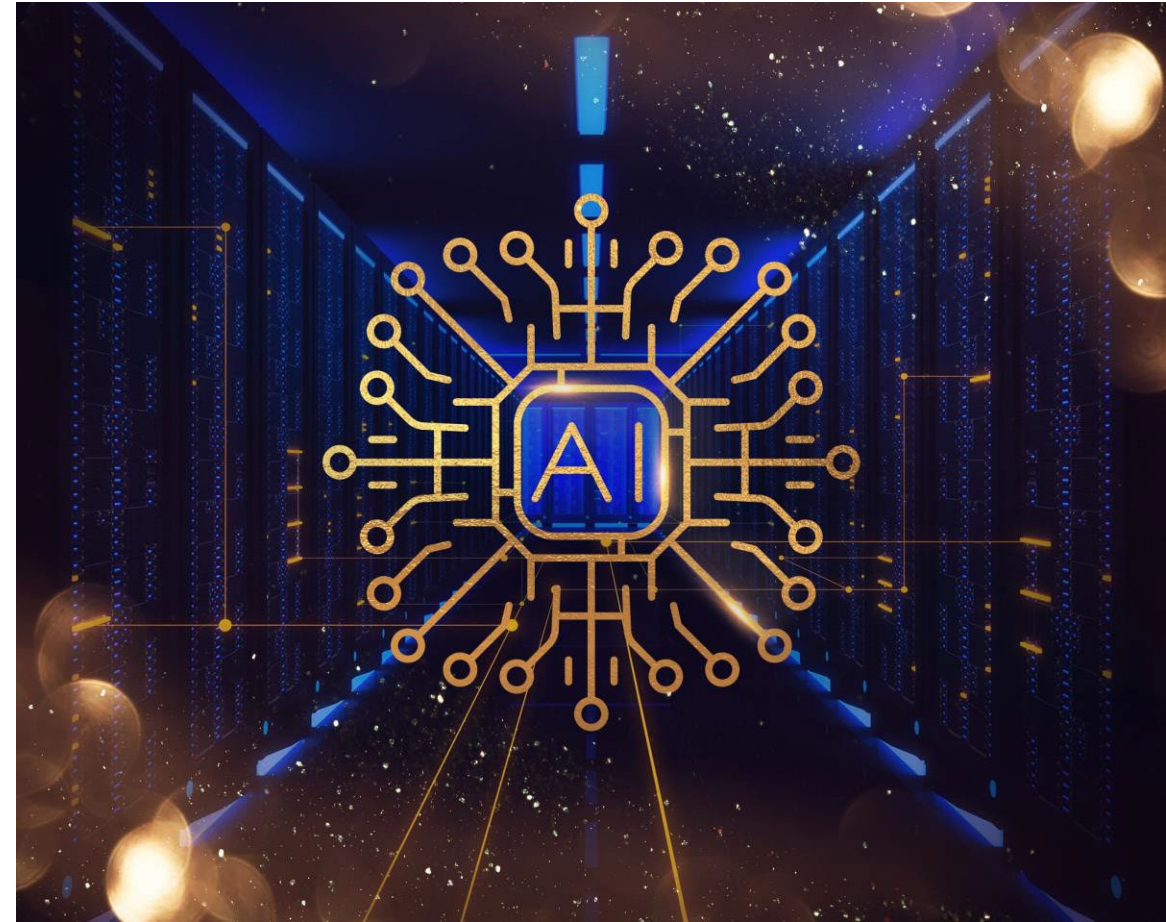


.
.
.
.



Overview


- Basics of Recommendation Engines
- Types of Recommender
- Collaborative Filtering
- *K*-Nearest Neighbour
- Content-based Filtering
- Predictive Analytics as Recommender
- Prescriptive Analytics as Recommender



.
.

At the end of this you should be able to

- Understand about the recommendation engine
- Understand about the data for the recommendation engine
- Understand different recommendation system
- Understand applications of recommendation system






Recommender Systems

Recommender/Recommendation Engine

- A recommendation engine, a recommender, is an AI system that suggests items to a user. Recommendation systems rely on data-driven analytics and ML algorithms to find patterns in user behaviour data and recommend relevant items based on those patterns.
- Recommendation engines help users discover content, products or services they might not have found on their own
- The suggestions created by recommendation systems also play a vital role in personalising user experiences.
- **Examples:** video to watch, a similar song to listen to, relevant search results or a product that complements a particular order.

Recommendation Engine examples

- **Facebook**—“People You May Know”
- **Netflix**—“Other Movies You May Enjoy”
- **LinkedIn**—“Jobs You May Be Interested In”
- **Amazon**—“Customer who bought this item also bought ...”
- **YouTube**—“Recommended Videos”
- **Google**—“Search results adjusted”
- **Pinterest**—“Recommended Images”



How Recommendation Engine works

Collect Data: User activities, products, demographics, psychographics (interests or lifestyle)

Data Storage:

- A data warehouse can aggregate data from different sources to support data analysis and machine learning.
- A data lakehouse combines the best aspects of data warehouses and data lakes into a single data management solution

Analysis: ML algorithms trained on large datasets detect patterns, identify correlations and weigh the strength of those patterns and correlations.

Filtering:

- Showing the most relevant items
- Apply specific mathematical rules and formulas to the data depending on the type of recommendation engine used.

Refining: Regularly assess the outputs of a recommendation system and further optimise the model

Types of Recommendation Engines

Collaborative Filtering

- Filter suggestions are based on a particular user's likeness to others.
- Assume that users with comparable preferences will likely be interested in the same items
- **Example:** Amazon product recommendations, Spotify recommendations
- **Limitation:** cold start problem, which happens when the system has limited historical data to draw from, especially for new users.

Content-based Filtering

- Filters recommendations based on an item's features.
- Assume that if a user likes a particular item, they will also like another similar item
- **Example:** Recommending a movie with similar genres, actors, or directors to movies a user has previously watched and enjoyed

Hybrid Recommendation System

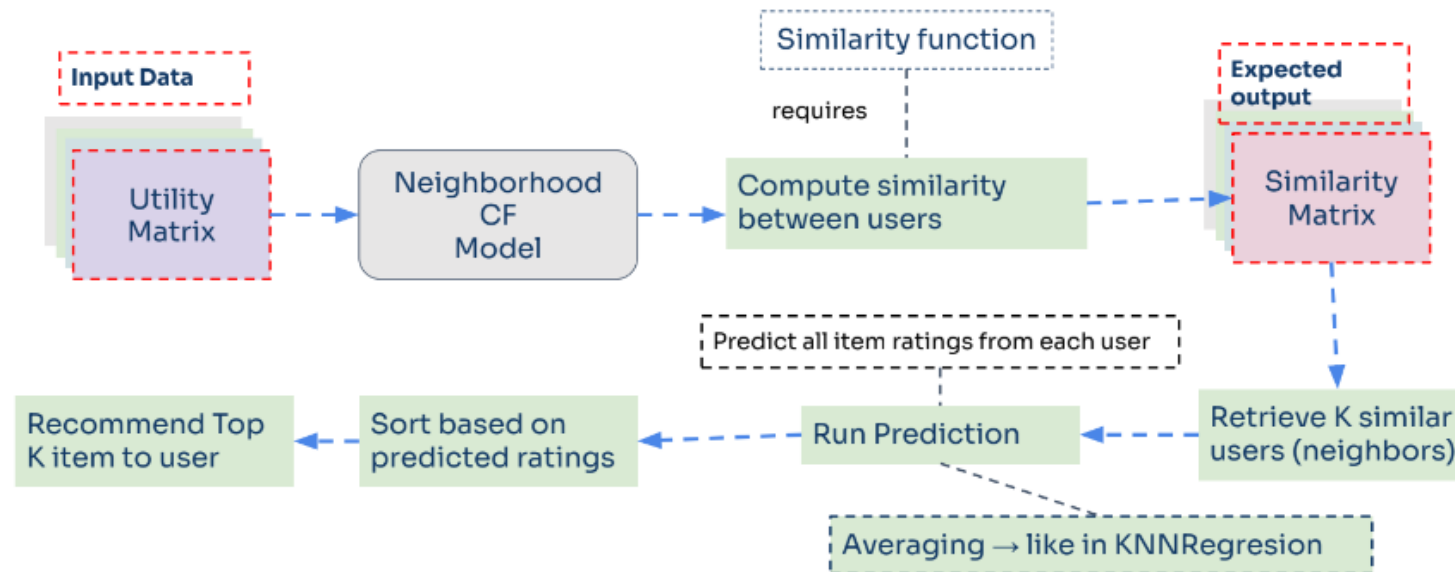
- Merges collaborative filtering and content-based filtering.
- A hybrid approach can greatly enhance a recommendation engine's performance but requires advanced architectures and intensive computational power.
- **Example:** Netflix

Collaborative Filtering

Memory-based

Represent users and items as a matrix. They are an extension of the k -nearest neighbours (kNN) algorithm because they aim to find their “nearest neighbours,” which can be similar users or items. There are two types:

- ✓ User-based filtering computes similarities between a particular user and all other users in the matrix
- ✓ Item-based filtering computes item similarity through user behaviour (how users interact with items, not item features). No **ML models** are used.



Collaborative Filtering

Model-based

- Create a predictive machine learning model of the data.
- The user-item matrix serves as the training data set for the model, which then yields predictions for missing values, that is, items that a user has not yet found and will, therefore, be recommended.
- Use matrix factorisation
- A more advanced implementation of matrix factorisation harnesses deep learning neural networks. Other model-based systems employ machine learning algorithms such as Bayes classifiers, clustering and decision trees.

Collaborative Filtering

Item - What the system recommends to the user (CD, news, books, movies...)

User preferences - ratings for products

User actions - user browsing history


The task of a CF algorithm is to find the item likeness of two forms :

1. Prediction – a numerical value expressing the predicted likelihood of an item for the active user
2. Recommendation – a list of N items that the active user will like the most

k - Nearest Neighbour Algorithm

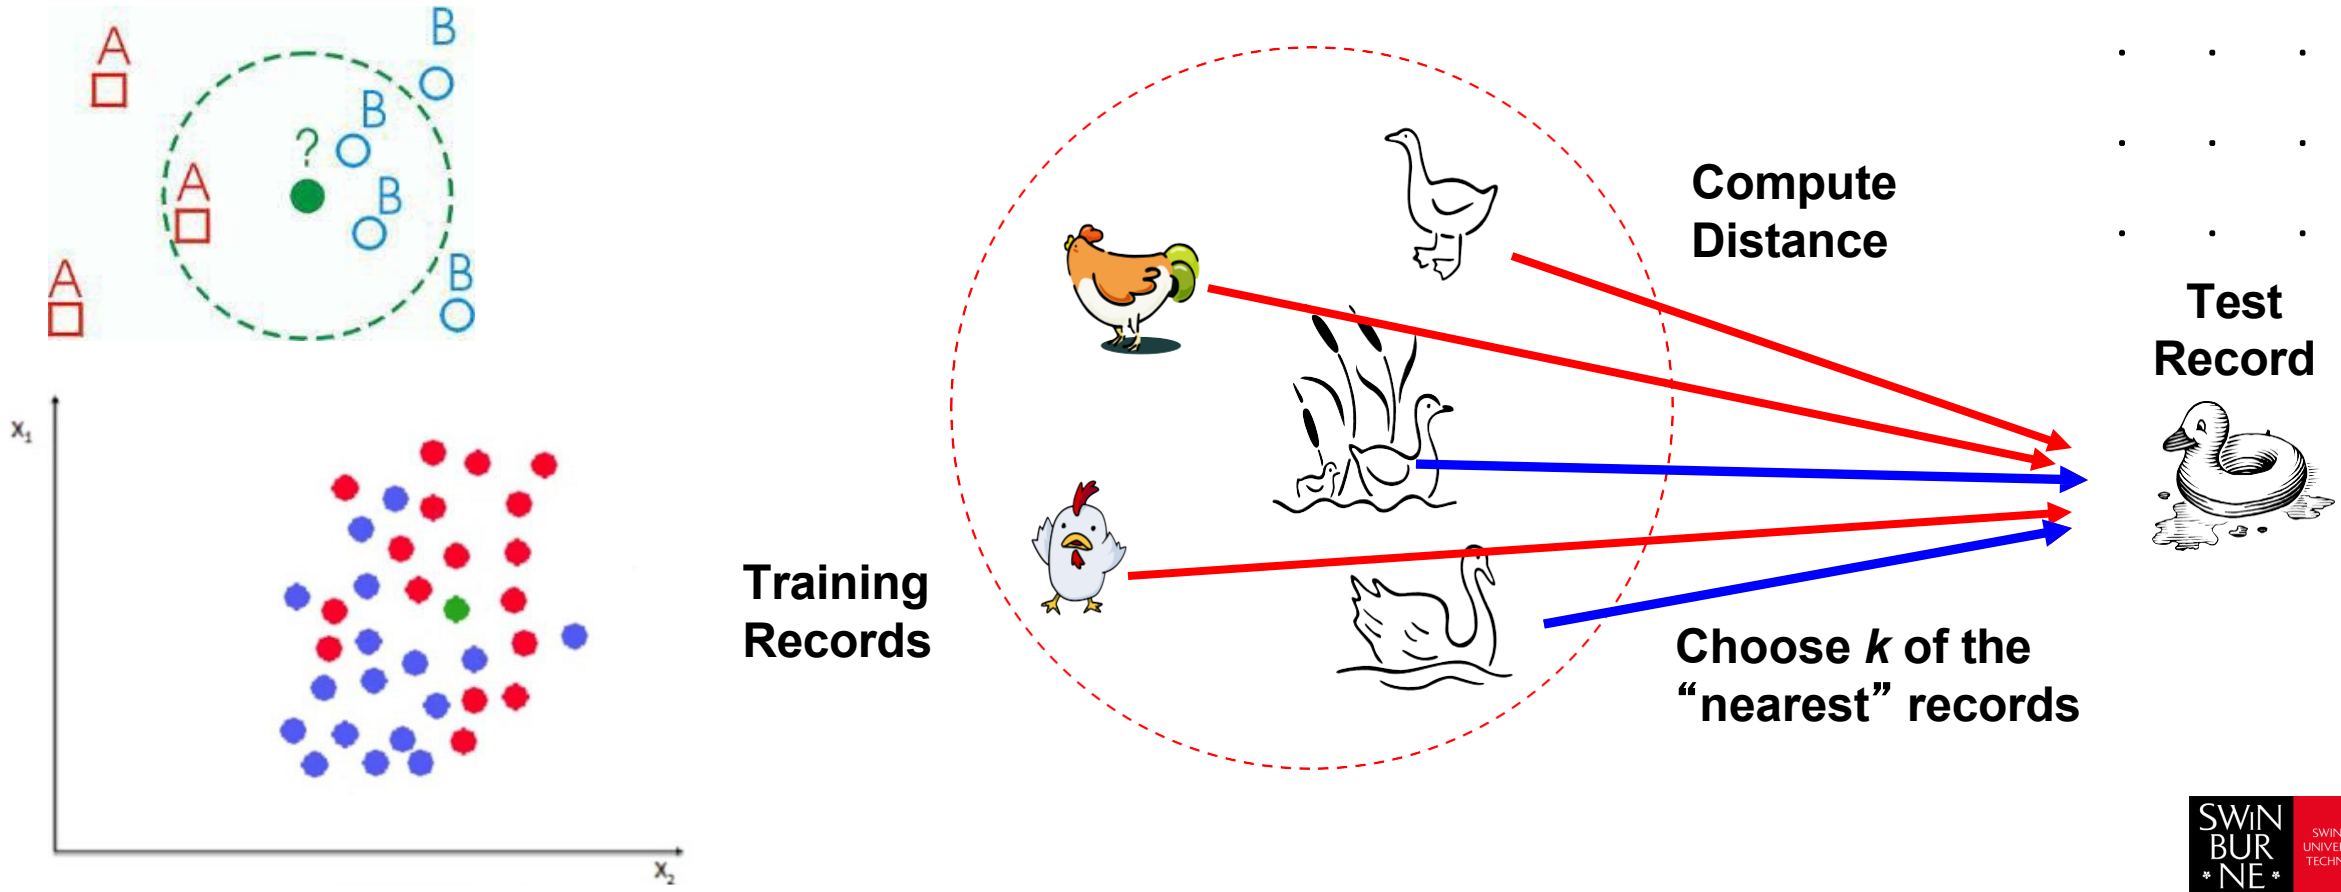
Tell me about your friends(*who your neighbours are*), and *I will tell you who you are*.

- A distance measure is needed to determine the “closeness” of instances
 - Classify an instance by finding its nearest neighbours and picking the most popular class among the neighbours
-
- Simple to implement and use
 - Comprehensible – easy to explain the prediction
 - Robust to noisy data by averaging k -nearest neighbours



KNN approach


- An object (a new instance) is classified by a majority vote for its neighbour classes.
- The object is assigned to the most common class amongst its K nearest neighbours (*measured by a distance function*)



Distance between neighbours


- Calculate the distance between the new example (E) and all examples in the training set.
- *Euclidean* distance between two examples,
 - $X = [x_1, x_2, x_3, \dots, x_n]$
 - $Y = [y_1, y_2, y_3, \dots, y_n]$
 - The Euclidean distance between X and Y is defined as:

$$D(X, Y) = \sqrt{\sum_{k=1}^n (x_k - y_k)^2}$$



K - Nearest Neighbour Algorithm


- All the instances correspond to points in an n -dimensional feature space.
- Each instance is represented with a set of numerical attributes.
- Each training data consists of vectors and a class label associated with each vector.
- Classification is done by comparing feature vectors of different K nearest points.
- Select the K -nearest examples to E in the training set.
- Assign E to the most common class among its K -nearest neighbours.




How to choose K

If K is too small, it is sensitive to noise points.


Larger K works well. But too large K may include majority points from other classes.



The rule of thumb is $K < \sqrt{n}$, and n is a number of examples.




How to choose K



(a) 1-nearest neighbour (b) 2-nearest neighbour (c) 3-nearest neighbour

K -nearest neighbours of a record x are data points that have the k smallest distances to x




KNN Feature Weighting

- Scale each feature by its importance for classification,


$$D(a, b) = \sqrt{\sum_k w_k (a_k - b_k)^2}$$

- Can use our prior knowledge about which features are more important
- Can learn the weights w_k using **cross-validation** (to be covered later)




KNN Feature Normalisation

- The distance between neighbours could be dominated by some attributes with relatively large numbers, e.g., the income of customers in our previous example.
- Arises when two features are on different scales.
- It is important to normalise those features.
 - Mapping values to numbers between 0 and 1.



KNN Classification


Age	Loan	Default	Distance
0.125	0.11	N	0.7652
0.375	0.21	N	0.5200
0.625	0.31	N	0.3160
0	0.01	N	0.9245
0.375	0.50	N	0.3428
0.8	0.00	N	0.6220
0.075	0.38	Y	0.6669
0.5	0.22	Y	0.4437
1	0.41	Y	0.3650
0.7	1.00	Y	0.3861
0.325	0.65	Y	0.3771
0.7	0.61	?	



$$X_s = \frac{X - Min}{Max - Min}$$

Collaborative Filtering Steps

- How do you determine which users or items are similar to one another?
- Given that you know which users are similar, how do you determine the rating that a user would give to an item based on the ratings of similar users?
- How do you measure the accuracy of the ratings you calculate?



Find Similar Users on the Basis of Ratings


- Users A, B, C, and D, who have rated two movies

Ratings by A are [1.0, 2.0].

Ratings by B are [2.0, 4.0].

Ratings by C are [2.5, 4.0].

Ratings by D are [4.5, 5.0].



Compute Similarity

```
from scipy import spatial
```

```
a = [1, 2]  
b = [2, 4]  
c = [2.5, 4]  
d = [4.5, 5]
```

```
spatial.distance.euclidean(c, a)  
2.5
```

```
spatial.distance.euclidean(c, b)  
0.5
```

```
spatial.distance.euclidean(c, d)  
2.23606797749979
```



Compute Rating

After determining a list of users similar to user U , you need to calculate the rating R that U would give to a certain item I ,

$$R_U = \frac{\sum_{u=1}^n R_u}{n}$$



Collaborative Filtering

- **User-based:** For a user U , with a set of similar users determined based on rating vectors consisting of given item ratings, the rating for an item I , which hasn't been rated, is found by picking out N users from the similarity list who have rated the item I and calculating the rating based on these N ratings.
- **Item-based:** For an item I , with a set of similar items determined based on rating vectors consisting of received user ratings, the rating by a user U , who hasn't rated it, is found by picking out N items from the similarity list that U has rated and calculating the rating based on these N ratings.

Python Program for Collaborative Filtering

```
# load_data.py
import pandas as pd
from surprise import Dataset
from surprise import Reader

# This is the same data that was plotted for similarity earlier
# with one new user "E" who has rated only movie 1
ratings_dict = {
    "item": [1, 2, 1, 2, 1, 2, 1, 2, 1],
    "user": ['A', 'A', 'B', 'B', 'C', 'C', 'D', 'D', 'E'],
    "rating": [1, 2, 2, 4, 2.5, 4, 4.5, 5, 3],
}

df = pd.DataFrame(ratings_dict)
reader = Reader(rating_scale=(1, 5))

# Loads Pandas dataframe
data = Dataset.load_from_df(df[["user", "item", "rating"]], reader)

# Loads the built-in Movielens-100k data
movielens = Dataset.load_builtin('ml-100k')
```

Using KNN

```
# recommender.py
from surprise import KNNWithMeans

# To use item-based cosine similarity
sim_options = {
    "name": "cosine",
    "user_based": False, # Compute similarities between items
}

algo = KNNWithMeans(sim_options=sim_options)

from load_data import data
from recommender import algo
trainingSet = data.build_full_trainset()
algo.fit(trainingSet)

# Computing the cosine similarity matrix...
# Done computing similarity matrix.
# <surprise.prediction_algorithms.knns.KNNWithMeans object at 0x7f04fec56898>
prediction = algo.predict('E', 2)
>>> prediction.est 4.15
```

Content-Based Filtering

Benefits:

- Independent of other user data
- Tailored to user preferences - aligning recommendations with the user's interests and preferences.
- Transparency in recommendations is directly tied to the user's actions.
- Simplicity in creation and data science- focus primarily on classifying items based on attributes, leveraging techniques such as vector space models and term frequency analysis.
- Overcoming the “cold start” problem -only initial user inputs deliver quality recommendations.

Challenges:

- Limited novelty and diversity may suggest overly familiar options and limit users’ diversity of possibilities.
- Scalability and attributes—Adding a new item, product, service, or content piece necessitates defining and tagging its attributes.
- Accuracy and attribute assignment - the precision and uniformity in assigning attributes significantly affect its success.
- Over-reliance on content quality and availability- highly dependent on item metadata or descriptions' quality, accuracy, and availability.

Difference Between Content-Based Filtering and Collaborative Filtering

Aspect	Content-Based Filtering	Collaborative Filtering
<i>Focus</i>	Item attributes	User Behaviour
<i>Recommendation</i>	Items similar to what the user likes	Items liked by similar users to the user
<i>Data required</i>	Information about the item	User behavior data, such as ratings or purchases
<i>Advantage</i>	Doesn't require user data	Can recommend niche or new items
<i>Disadvantage</i>	May miss out on new interests	Needs sufficient user data to be effective

Prescriptive Analytics-based Recommendation

- **Descriptive analytics** - Answers what has already happened
- **Diagnostic analytics** - Answers why did this happen
- **Predictive analytics** - Answers what could/might happen
- **Prescriptive analytics** - Answers what should happen/what should we do next

Prescriptive analytics is analysing data to identify patterns, which can be used to make predictions and determine optimal courses of action.

- Its focus is predicting future outcomes and recommending actions or decisions to achieve desired outcomes or prevent undesirable ones.
- Prescriptive analytics adds a recommendation layer on top of predictive analytics.
- **Example:** An e-commerce platform can use prescriptive analytics to analyse customers' actions, such as browsing and buying habits, to give the most spot-on product recommendations.
- **Example:** Using prescriptive analytics to dive into transaction patterns and customer behaviour, they can spot fraud and leverage insights to implement more effective protection measures.

How Prescriptive Analytics-based Recommendation Works

How it works:

1. Data Collection and Analysis:

Prescriptive analytics uses various data sources to analyse patterns, trends, and relationships. . . .

2. Model Building:

Mathematical models and algorithms are used to simulate different scenarios and evaluate potential outcomes. . . .

3. Scenario Simulation:

The models explore various possibilities and consider constraints and objectives to determine the most effective course of action.

4. Recommendation Generation:

Based on the analysis and simulation, prescriptive analytics provides specific recommendations for decision-makers.

Learn, Practice and Enjoy the AI journey

